

What it costs to run — five line items

Serper, Apify, Amplify, the backend, and the LLM. Call volumes measured from the live pipeline; unit prices from each vendor.

3 Aug 2026 Basis measured, per 1,000 talent Region us-east-1

THE WHOLE PICTURE

Cost per 1,000 talent processed

One "talent" = one person checked across all five platforms. Fixed costs shown separately below.		
LINE ITEM	WHAT DRIVES THE COST	PER 1,000 TALENT
1. Serper EXTERNAL	10,880 searches — 5 per talent to find profiles, 5.9 to pull evidence \$0.75 per 1,000 searches (mid-tier)	\$8.16
2. Apify EXTERNAL	Runs only for talent with an unresolved platform — ~875 of 1,000 \$0.002 per result + \$0.001 per actor start	\$2.10 – \$7.35
3. Amplify AWS	Frontend hosting — fixed, does not scale with rows Builds, bandwidth, SSR compute	\$15 / mo
4. Backend AWS	App Runner + S3 + Secrets Manager + CloudWatch — fixed See recommendation below	\$30 / mo
5. LLM AWS	4,880 calls · 11.6 M input + 1.46 M output tokens Claude Sonnet 5 with prompt caching	\$40.34
Variable subtotal, per 1,000 talent		\$50.60 – \$55.85
Fixed, per month		\$45

Backend — what I'd use, and why

AWS App Runner. It takes your existing FastAPI container and runs it with no load balancer, no VPC, no NAT gateway, and no cluster to manage. That removes about \$120/month of plumbing you would otherwise pay for on ECS Fargate, and roughly a week of setup.

SERVICE	CONFIGURATION	\$ / MONTH
App Runner	1 vCPU / 2 GB — memory billed always, CPU only while working	\$24
S3	Uploaded files and generated exports	\$1
Secrets Manager	API keys for the four services above	\$2
CloudWatch	Logs and alarms	\$3
Backend total		\$30

THE TRADE-OFF YOU'RE ACCEPTING

App Runner can scale an instance in while a long run is still going, and the job dies with it — results are currently held in memory, so that run is lost. Fine for a pilot; before production you want jobs in a database and the pipeline in a separate worker. That's the upgrade to ECS Fargate and roughly **+\$130/month**, worth deferring until volume justifies it.

LLM — the biggest lever

Same 11.6 M input / 1.46 M output tokens per 1,000 talent, via Claude Platform on AWS.			
MODEL	RATE IN / OUT	PER 1,000 TALENT	NOTE
Claude Haiku 4.5	\$1 / \$5	\$18.90	Cheapest, but accuracy unproven on this task
Claude Sonnet 5	\$3 / \$15	\$40.34	Recommended. \$27.20 at the intro rate through 31 Aug
Claude Opus 5	\$5 / \$25	\$67.24	Overkill for the whole run

FREE 29% SAVING, ALREADY INCLUDED

The system prompt is 1,244 tokens and identical on every call, so prompt caching bills half the input at one-tenth rate. Without it Sonnet 5 costs \$56.70 instead of \$40.34. Note this does **not** work on Haiku — its minimum cacheable prompt is 4,096 tokens.

RUN RATE

What you'd actually pay

Sonnet 5 with caching, Serper at \$0.75/1k, Apify mid-estimate.					
VOLUME	SERPER	APIFY	LLM	AMPLIFY + BACKEND	TOTAL / MONTH
5,000 talent / mo	\$41	\$24	\$202	\$45	\$312
25,000 talent / mo	\$204	\$118	\$1,009	\$55	\$1,386
100,000 talent / mo	\$816	\$473	\$4,034	\$175	\$5,498

The LLM is 65–73% of every scenario. If cost becomes a problem, that is the only line worth attacking — and the way to attack it safely is to build the accuracy test set first, then check whether Haiku holds up.

Three things to confirm

ITEM	WHY IT'S UNCERTAIN
Serper tier	Priced here at \$0.75 per 1,000 searches. Entry pricing is \$1.00, volume pricing nearer \$0.30 — a 3× spread on this line. Confirm the plan.
Apify results per talent	Unit prices are confirmed from the actor itself; how many dataset items it returns per name is not yet measured, hence the range. One pilot run settles it.
Which LLM	Accuracy has not been measured on a labelled set, so the Haiku saving — 53% of the largest line — can't be taken safely yet.

Call volumes measured by instrumenting the pipeline against the production brand file on 3 Aug 2026. Apify unit prices read from the actor's own pricing record via the Apify API. AWS and LLM rates are published list prices and should be re-checked at purchase.